ReviewWise — M3 summary draft

Team: MAKNI Abdourrahmane · HAFT Benjamin · ABDAI Eman

Deadline: 16 November 2026, 23:59.

Repository: https://github.com/abdai-eman/reviewwise-ai

M3 tagged commit and unlisted video: pending — not created or published.

Problem and users. Independent café owners in Luxembourg receive multilingual feedback and need to decide what to improve. ReviewWise aims to turn supplied reviews into inspectable complaints and practical actions.

Product. A local Streamlit page accepts pasted reviews or a UTF-8 CSV. Ollama with `qwen2.5:3b` extracts feedback and original-language evidence. Python checks JSON, review/category/evidence IDs and quotations, counts distinct review IDs per complaint, and creates a downloadable weekly report with up to three supported actions. Current source groups up to four short reviews, shows live partial results, lists validation failures and supports stopping. Real-model verification of this correction is pending; the Windows ZIP still contains the earlier version. Up to 1,000 reviews are accepted. The Monday–Sunday week is selected by the user; input dates are not verified or filtered. All accepted reviews are counted. No database, cloud AI, API key or Docker is required. See `docs/architecture.svg`.

Installation. On Windows 11 with App Installer/WinGet: download the M3 ZIP, extract it, double-click `START_REVIEWWISE.bat`. First launch installs missing Python/Ollama only when missing, reuses Python 3.11-3.15, creates `.venv`, installs hash-checked packages, downloads the model and opens the app. Later starts use the same file. It asks before installing another Python. Without WinGet, installing missing prerequisites adds up to two steps: five steps total. The model is about 1.9 GB (unverified estimate); 16 GB RAM is a planning recommendation, not a measured minimum.

Evidence and status. 67 offline tests pass on Python 3.12, including 80-review handling, live partial coverage, session reuse, failure recovery and date labels. Earlier launcher checks covered Python 3.11-3.15; the latest detailed results are in `docs/m3_validation.md`. Package install and native Ollama 0.40.14 startup work. Cloud policy blocks the model registry, so no real M3 inference or accuracy improvement is claimed. M2 evidence is unchanged. Windows setup, a real-input run, another person's laptop test, offline operation and the video remain pending.

Known limitations. AI may miss or misclassify multilingual feedback, especially Luxembourgish; Python counts cannot correct incorrect extraction. Six café categories are limited and “other” combines different topics. CPU inference may be slow for a new weekly file. Validated analyses are reused in the current session; closing or restarting the app can lose them. Quotes establish traceability, not semantic truth. M3 performance/user benefit remain unmeasured.

Cost and privacy. 80 short new reviews can use 20 group calls, or up to 160 with long inputs/fallback × €0 metered API fee = €0 metered API fees. Hardware/electricity remain unmeasured. Reviews go only to local Ollama; no automatic report files are saved. Remove unnecessary personal information and use permitted data.

Who did what since M2. Actual member contributions require team confirmation. Codex assisted with M3 code, tests and documentation; that does not establish an individual student's contribution or technical understanding.

Before submission. Complete the real-input and other-laptop tests, record an unlisted three-minute video, replace pending links/contributions, review this draft, export the final one-page PDF and tag the agreed M3 commit.